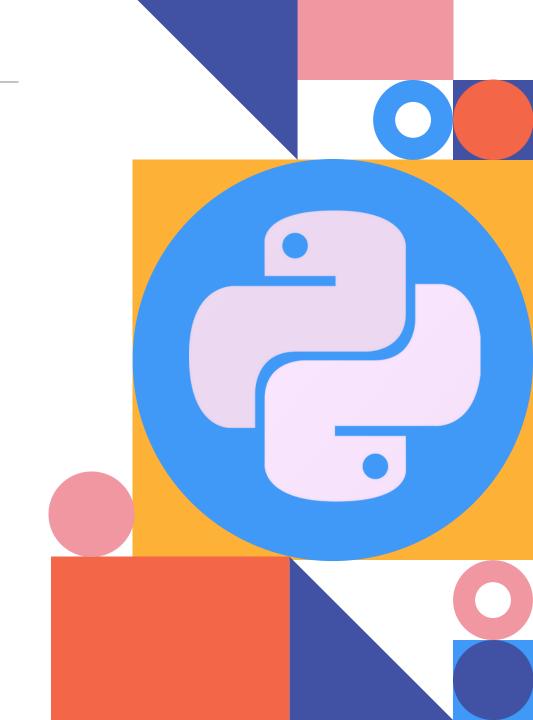
# Data Analysis II

Yan-Fu Kuo

Dept. of Biomechatronics Engineering National Taiwan University





#### Contents

01 Exploratory Data Analysis

02 Data Analysis Packages

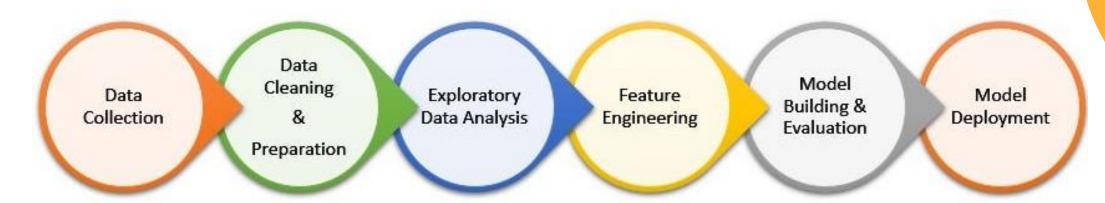
**Pandas** 

## 01 Exploratory Data Analysis

## Exploratory Data Analysis (EDA)

An approach to analyzing data sets to summarize their main characteristics, often with visual methods


- 1. Understand your variables
- 2. Clean your dataset
- 3. Analyze relationships between variables



#### COVID-19 in Taiwan



#### COVID-19 in Taiwan





### 02 Pandas

#### Pandas



A fast and efficient **DataFrame** object for data manipulation with integrated indexing

Tools for **reading and writing data** between inmemory data structures and different formats: CSV and text files, Microsoft Excel, SQL databases, and the fast HDF5 format

..

## Pandas Objects



#### Series

One-dimensional array of indexed data, or a specialization of a Python dictionary



#### Dataframe \*

A generalization of a NumPy array, or an extended Python dictionary




#### Index

An immutable array or as an ordered set



#### Series



What is the value of data[3]?

## Series as Specialized Dictionary

```
population_dict = {'California': 38332521,
                   'Texas': 26448193,
                   'New York': 19651127,
                   'Florida': 19552860,
                    'Illinois': 12882135}
population = pd.Series(population_dict)
population dict['California'] #38332521
population['California'] #38332521
population.sort_values(ascending=True)
```


38332521
16448193
19651127
19552860
12882135
12882135
16448193
19552860
19651127

38332521

California

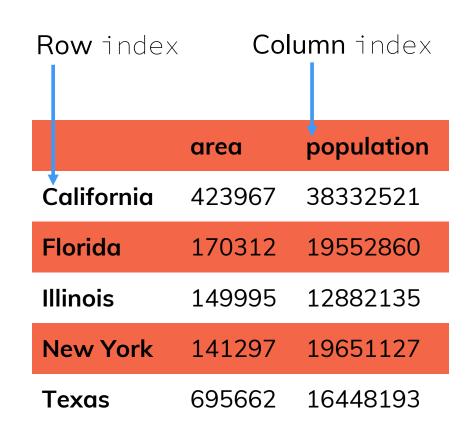
### Creating a DataFrame

- From a single Series object
- From a list of dictionaries
- From clipboard
- From a two-dimensional NumPy array
- From a NumPy structured array



#### DataFrame from Series

```
States['area']
```

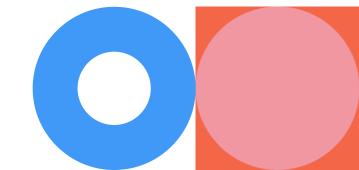


#### From a List of Dictionaries

```
dict = {
    "col 1": [1, 2, 3],
    "col 2": [10, 20, 30],
    "col 3": list('xyz'),
    "col 4": ['a', 'b', 'c'],
    "col 5": pd.Series(range(3))
}

df = pd.DataFrame(dict)
```

	col 1	col 2	col 3	col 4	col 5
0	1	10	X	а	0
1	2	20	У	b	1
2	3	30	Z	С	2



## From Clipboard

#### Copy the text:

```
a b c d
0 1 inf 1/1/00
2 71.38 N/A 5-Jan-13
4 54.59 nan 7/24/18
6 40.42 None NaT
```

```
pd.read_clipboard()
```

#### From CSV

```
df = pd.read_csv('titanic.csv')
df.head(6)
```

#### Also check:

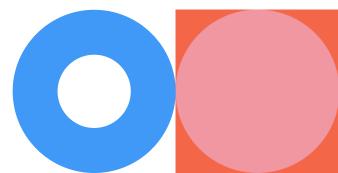
df.tail(6)

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th	female	38.0	1	0	PC 17599	71.2833	C85	С
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
5	6	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.4583	NaN	Q

#### Titanic Machine Learning Competition







### Titanic Dataset

	Passengerld	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr.	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John	female	38.0	1	0	PC 17599	71.2833	C85	С
2	3	1	3	Heikkinen,	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr.	male	35.0	0	0	373450	8.0500	NaN	S
888	889	0	3	Johnston, Miss. "	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr.	male	26.0	0	0	111369	30.0000	C148	С
890	891	0	3	Dooley, Mr. k	male	32.0	0	0	370376	7.7500	NaN	Q

## df = read\_csv()

pandas.read\_csv(filepath\_or\_buffer, sep=',', delimiter=None, header='infer', names=None, index col=None, usecols=None, squeeze=False, prefix=None, mangle\_dupe\_cols=True, dtype=None, engine=None, converters=None, true values=None, false values=None, skipinitialspace=False, skiprows=None, skipfooter=0, nrows=None, na\_values=None, keep\_default\_na=True, na\_filter=True, verbose=False, skip\_blank\_lines=True, parse\_dates=False, infer datetime format=False, keep date col=False, date parser=None, dayfirst=False, cache\_dates=True, iterator=False, chunksize=None, compression='infer', thousands=None, decimal='.', lineterminator=None, quotechar='"', quoting=0, doublequote=True, escapechar=None, comment=None, encoding=None, dialect=None, error bad lines=True, warn bad lines=True, delim whitespace=False, low\_memory=True, memory\_map=False, float precision=None)

Memorize? Try SHIFT + TAB in jupyter

## df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#
      Column
                        Non-Null Count
                                          Dtype
      PassengerId
                        891 non-null
                                          int64
0
      Survived
                        891 non-null
                                          int64
      Pclass
                        891 non-null
                                          int64
                                          object
      Name
                       891 non-null
                                          object
      Sex
                        891 non-null
      Age
                        714 non-null
                                          float64
      SibSp
                        891 non-null
                                          int64
                                          int64
      Parch
                        891 non-null
                                          object
      Ticket
                        891 non-null
      Fare
                        891 non-null
                                          float64
10
                        204 non-null
                                          object
      Cabin
11
      Embarked
                        889 non-null
                                          object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

#### Q: What are types of the objects?


df.info(memory\_usage="deep")

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
     Column
                       Non-Null Count
                                         Dtype
0
      PassengerId
                       891 non-null
                                         int64
     Survived
                       891 non-null
                                         int64
      Pclass
                       891 non-null
                                         int64
                                         object
     Name
                       891 non-null
      Sex
                       891 non-null
                                         object
     Age
                       714 non-null
                                         float64
     SibSp
                       891 non-null
                                         int64
     Parch
                       891 non-null
                                         int64
     Ticket
8
                       891 non-null
                                         object
     Fare
                       891 non-null
                                         float64
     Cabin
                       204 non-null
                                         object
10
      Embarked
                       889 non-null
                                         object
11
dtypes: float64(2), int64(5), object(5)
memory usage: 315.0 KB
```

## df.describe()

	Passengerld	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

## Dealing with Big Data



## Titanic Dataset (Review)

	Passengerld	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr.	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John	female	38.0	1	0	PC 17599	71.2833	C85	С
2	3	1	3	Heikkinen,	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr.	male	35.0	0	0	373450	8.0500	NaN	S
•••	•••	•••			•••				•••	•••		
888	889	0	3	Johnston, Miss. "	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr.	male	26.0	0	0	111369	30.0000	C148	С
890	891	0	3	Dooley, Mr. k	male	32.0	0	0	370376	7.7500	NaN	Q

## Reducing Memory Usage – 1

Using the Pandas Category data type

```
10 Cabin 204 non-null
11 Embarked 889 non-null
dtypes: float64(2), int64(5), object(5)
memory usage: 315.0 KB
```

```
10 Cabin 204 non-null
11 Embarked 889 non-null
dtypes: float64(2), int64(5), object(5)
memory usage: 135.0 KB
```



## Reducing Memory Usage – 2

**Batch processing** 

```
reader = pd.read_csv('titanic.csv', chunksize=4)
for _, df_partial in zip(range(2), reader):
    display(df_partial)
```

Q: How do you determine the batch size?



## Customizing DataFrame Display

```
print("Default display.max_colwidth:",
    pd.get_option("display.max_colwidth"))
```

Default display.max colwidth: 10

```
pd.set_option("display.max_colwidth", 5)
```

## Customizing DataFrame Display

```
pd.set_option("display.max_columns", 5)
pd.set_option("display.max_rows", 6)
```

	Passengerld	Survived	 Cabin	Embarked
0	1	0	 NaN	S
1	2	1	 C85	С
2	3	1	 NaN	S
	•••		 •••	
888	889	0	 NaN	S
889	890	1	 C148	С
890	891	0	 NaN	Q



## Exercise (5 mins)

	Passengerld	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr	male	22.0	1	0	A/5 21171	7.25	NaN	S
1	2	1	1	Cumings, Mr	female	38.0	1	0	PC 17599	71.28	C85	С
2	3	1	3	Heikkinen,	female	26.0	0	0	STON/O2. 31	7.92	NaN	S
3	4	1	1	Futrelle, M	female	35.0	1	0	113803	53.10	C123	S
•••					•••						•••	
887	888	1	1	Graham, Mis	female	19.0	0	0	112053	30.00	B42	S
888	889	0	3	Johnston, M	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	890	1	1	Behr, Mr. K	male	26.0	0	0	111369	30.00	C148	С
890	891	0	3	Dooley, Mr	male	32.0	0	0	370376	7.75	NaN	Q

## DataFrame Styling

Colorful Titanic Dataset

Passengerld	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
543	0	3	female	11	4	2	347082	31.3	nan	S
243	0	2	male	29	0	0	W./C. 14263	10.5	nan	S
197	0	3	male	nan	0	0	368703	7.8	nan	Q
482	0	2	male	 nan	0	0	239854	0.0	nan	S
683	0	3	male	20	0	0	6563	9.2	nan	S
693	1	3	male	nan	0	0	1601	56.5	nan	S
510	1	3	male	26	0	0	1601	56.5	nan	S
511	1	3	male	29	0	0	382651	7.8	nan	Q
568	0	3	female	29	0	4	349909	21.1	nan	S
227	1	2	male	19	0	0	SW/PP 751	10.5	nan	S
722	0	3	male	17	1	0	350048	7.1	nan	S
645	1	3	female	1	2	1	2666	19.3	nan	С
681	0	3	female	nan	0	0	330935	8.1	nan	Q
114	0	3	female	20	1	0	4136	9.8	nan	S
630	0	3	male	nan	0	0	334912	7.7	nan	Q
411	0	3	male	nan	0	0	349222	7.9	nan	S
130	0	3	male	45	0	0	347061	7.0	nan	S

### Exercise – Data Cleansing (5 mins)

	Passengerld	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
542	543	0	3	female	11.0	4	2	347082	31.2750	NaN	S
242	243	0	2	male	29.0	0	0	W./C. 14263	10.5000	NaN	S
196	197	0	3	male	NaN	0	0	368703	7.7500	NaN	Q
481	482	0	2	male	NaN	0	0	239854	0.0000	NaN	S
682	683	0	3	male	20.0	0	0	6563	9.2250	NaN	S

```
Try:
```

df.fillna(₀)

df.fillna("Unknown")

df.Age.fillna(int(df.Age.mean()))

Q: Which on is better?

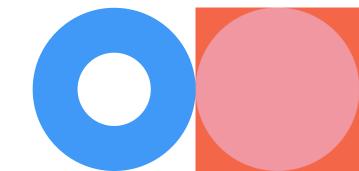
### Drop (Useless) Data

#### Drop column data

```
columns = ['Name', 'Ticket']
df.drop(columns, axis=1)
```

#### Drop row data

```
df.drop(2)
```



## DataFrame Indexing

df.Pclass

df[0:4]

0	3
1	1
2	3
3	1
887	1
888	3
889	1
890	3

	Passengerld	Survived		Embarked
0	1	0	•••	S
1	2	1	•••	С
2	3	1		S
3	4	1		S

- Index refers to columns and slicing refers to rows
- <u>Cannot</u> do indexing on both columns and rows together!

#### Indexer

Allow indexing DataFrame as if it is a simple NumPy array.

- df.loc is primarily label based, but may also be used with a boolean array.
- df.iloc is primarily integer position based (from 0 to length-1 of the axis), but may also be used with a boolean array.

#### Data Selection with Indexer

```
df.loc[0:4, 'PassengerId': 'Age']
```

```
df.iloc[0:5, 0:5]
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	female	38.0	1	0	PC 17599	71.2833	C85	С
2	3	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	male	35.0	0	0	373450	8.0500	NaN	S

#### Data Selection with Indexer

df.loc[df.Survived == 1, :]

	Passengerld	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	2	1	1	female	38.0	1	0	PC 17599	71.2833	C85	С
2	3	1	3	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	female	35.0	1	0	113803	53.1000	C123	S
8	9	1	3	female	27.0	0	2	347742	11.1333	NaN	S
9	10	1	2	female	14.0	1	0	237736	30.0708	NaN	С
875	876	1	3	female	15.0	0	0	2667	7.2250	NaN	С
879	880	1	1	female	56.0	0	1	11767	83.1583	C50	С


## Data Selection – String and Num

#### Masking

```
female_and_age_under_20 = (df.Sex == 'female') & (df.Age < 20)
df[female_and_age_under_20]</pre>
```

#### Query

```
age = 70
df.query("Age > @age & Sex == 'male'")
```



## Data Selection – String

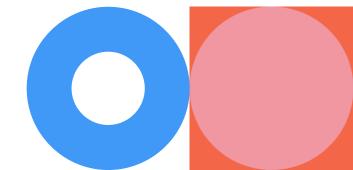
```
df[df.Name.str.contains("Mr\.")]
```

```
tickets = ["SC/Paris 2123", "PC 17475"]
df[df.Ticket.isin(tickets)]
```

```
top_k = 3
top_tickets = df.Ticket.value_counts()[:top_k]
top_tickets.index
```

### Exercise

Show the passenger info of the most frequently purchased class of tickets



## GroupBy: Split, Apply, Combine

df.groupby("Pclass").Age.mean()

					Pdass	Age						
				0	1	45		Pdass	Age(mean)			
	<b>Pdass</b>	Age	1	5	1	35	<b></b>	1	38.23	1		
0	1	45		<u> </u>			-					
1	2	20			Pdass	Age					<u>Pdass</u>	Age(mean)
								<b>Pdass</b>	Age(mean)		1	38.23
2	3	32	<b>/</b> →	1	2	20		2	29.87	<b>→</b>	2	29.87
3	2	26		3	2	26			23.07	1		
4	3	19					-				3	25.14
	3				Pdass	Age	_					
5	1	35		2	3	32		Pdass	Age(mean)			
			_	4	3	19	<b></b>	3	25.14	•		

## Insights

Does sex related to the survival rate in that accident?

```
df.groupby(['Sex','Survived']).Sex.count()
```

Sex	Survived				No Survived
female	0	81			Survived
	1	233			- Cal Tivea
male	0	468			
	1	109			
Name: S	ex, dtype:	int64	female	male	

#### Exercise

Does ticket class (pclass) related to the survival rate in that accident?

