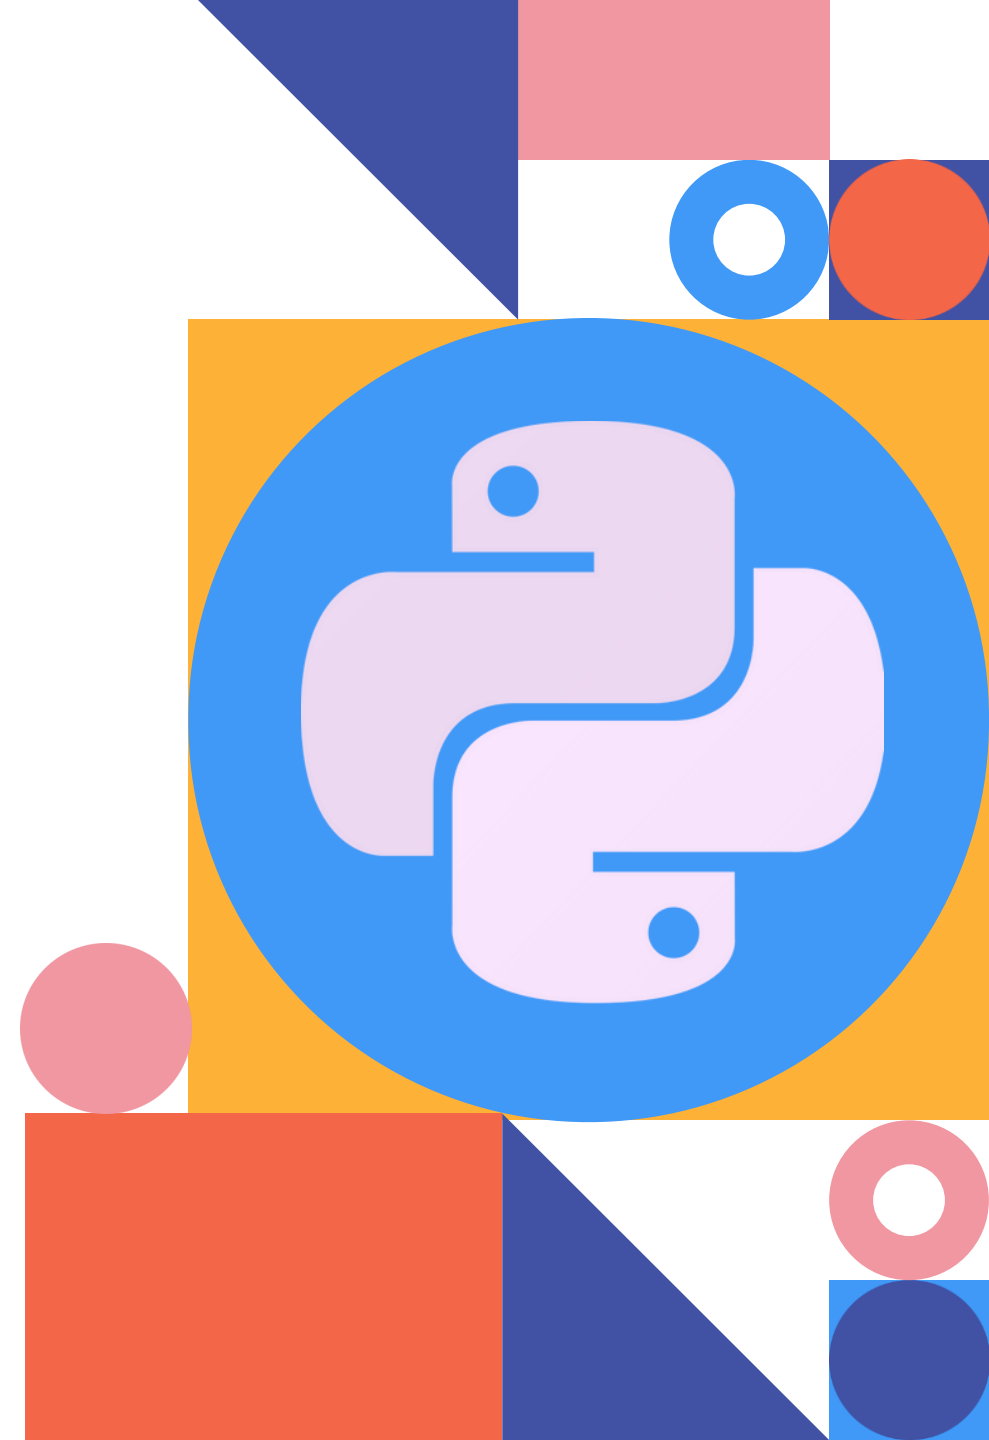

Deep Learning

Yan-Fu Kuo

Dept. of Biomechatronics Engineering
National Taiwan University

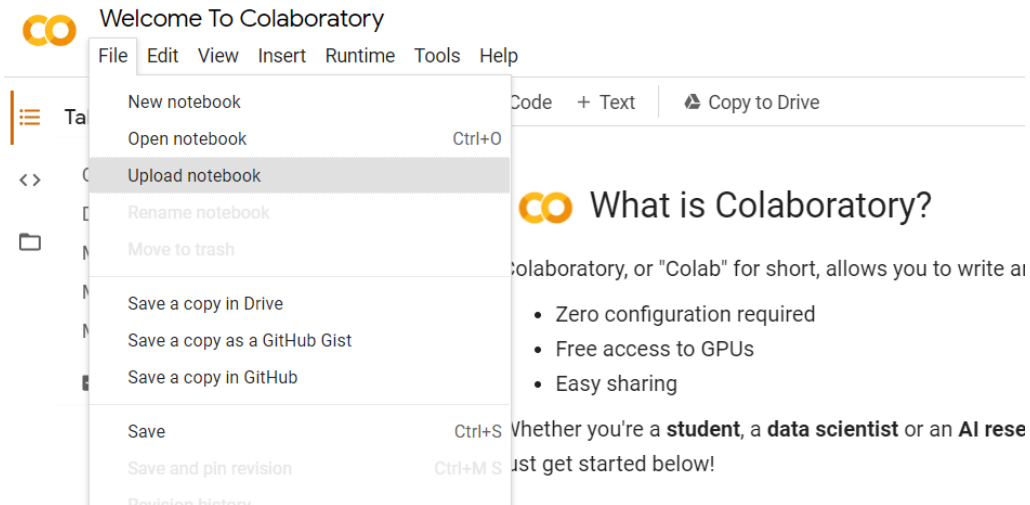


Google Colab

Colab is a hosted Jupyter notebook service that requires no setup to use, while providing free access to computing resources including GPUs.

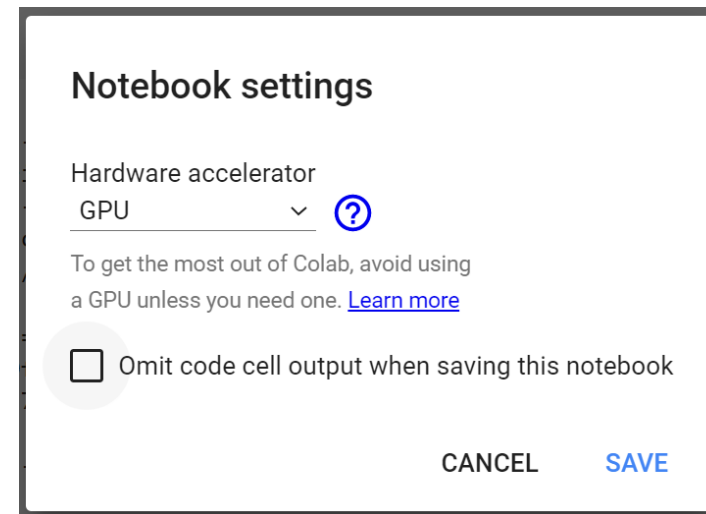
1

Upload 12DeepLearning.ipynb to Google Colab

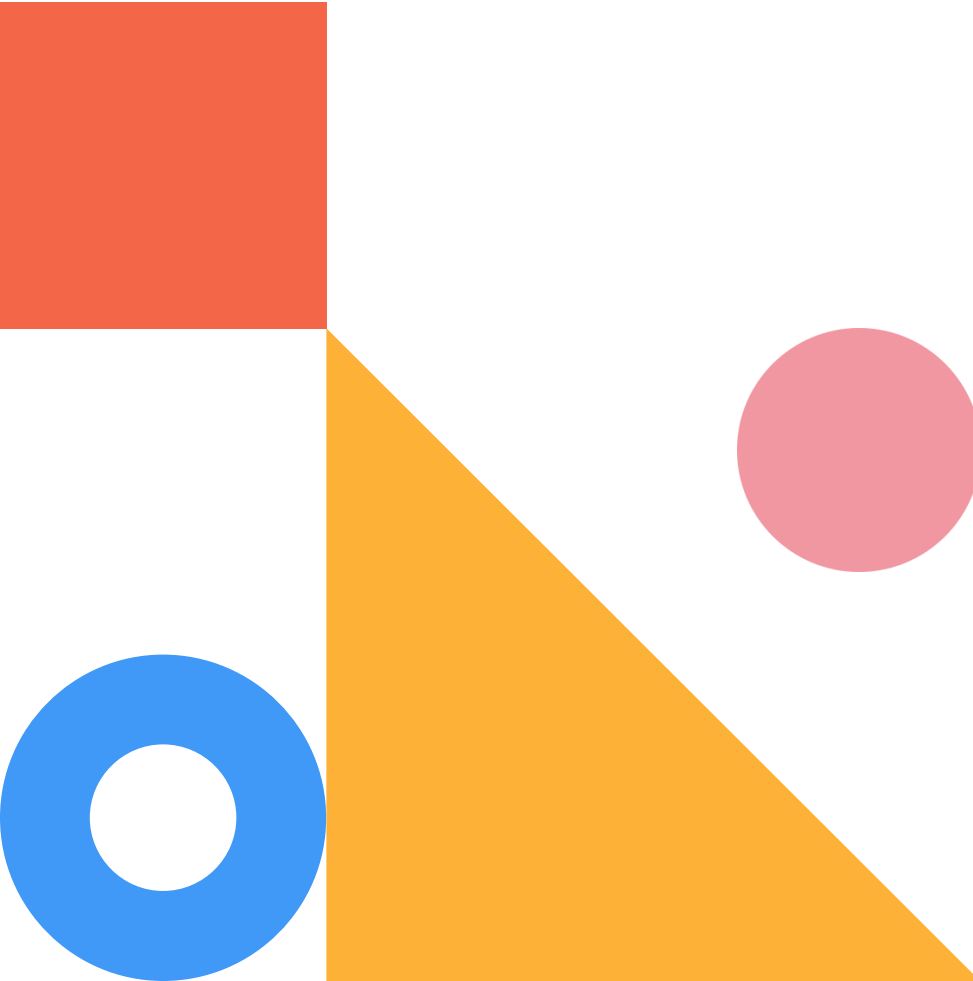
**2**

Setting Free GPU

Edit > Notebook settings > Change runtime type and select GPU as Hardware accelerator.



Contents



01 PyTorch and Tensor

02 Artificial Neural Networks

03 How to Train a Model

04 Convolutional Neural Networks

01 PyTorch and Tensor



PyTorch



PyTorch

A replacement for NumPy to use
the power of GPUs

A deep learning research platform
that provides maximum flexibility
and speed

Pytorch -> Get Started



GPU Acceleration



Tensors

Tensors are similar to NumPy's ndarrays, with the addition being that Tensors can also be used on a GPU to accelerate computing.

```
import torch
```


```
x = torch.tensor([5.5, 3])
```

```
tensor([5.5000, 3.0000])
```

```
import numpy as np
```

```
y = np.array([5.5, 3])
```

```
[5.5000, 3.0000]
```



Tensors | Operations

Arithmetic

```
x = torch.linspace(1, 10, 10)
y = torch.ones((1, 10))
```

```
x + y
```

```
tensor([[ 2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10., 11.]])
```

Slicing

```
x[:5]
```

```
tensor([1., 2., 3., 4., 5.])
```

Resizing

```
x = torch.randn(4, 4)
y = x.view(16)
z = x.view(-1, 8)
print(x.shape, y.shape, z.shape)
```

```
torch.Size([4, 4]) torch.Size([16]) torch.Size([2, 8])
```

.item()

```
x = torch.randn(1)
print(x.item())
```

```
-0.408441
```


Tensors | Numpy Bridge

NumPy Array → Torch Tensor


```
a = np.ones(5)
b = torch.from_numpy(a)
print(a)
print(b)
```

```
[1.  1.  1.  1.  1.]
tensor([1., 1., 1., 1., 1.],
        dtype=torch.float64)
```

Torch Tensor → NumPy Array

```
a = torch.ones(5)
print(a)
b = a.numpy()
print(b)
```

```
tensor([1., 1., 1., 1., 1.])
[1., 1., 1., 1., 1.]
```





Tensors | CUDA

```
x = torch.randn(1)
if torch.cuda.is_available():
    device = torch.device("cuda")
    y = torch.ones_like(x, device=device)
    x = x.to(device)
    z = x + y
    print(z)
    print(z.to("cpu", torch.double))
```

```
tensor([-0.5981], device='cuda:0')
tensor([-0.5981], dtype=torch.float64)
```

Tensors can be moved onto any device using the `.to` method.






02 Artificial Neural Networks

Deep Learning Models

Usually refers to neural networks with large numbers in layers and neurons.




Why Deep Learning?


Machine Learning




Deep Learning



Perceptron (Neuron)




$$y = \begin{cases} 1, & \Sigma > \text{threshold} \\ 0, & \Sigma < \text{threshold} \end{cases}$$




Activation Function

Step Function




$$f(x) = \begin{cases} 1, & x > 0 \\ 0, & x \leq 0 \end{cases}$$

Sigmoid




$$f(x) = \frac{1}{1 + e^{-x}}$$

ReLU




$$f(x) = \max(0, x)$$

TanH




$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


Linear



$$f(x) = x$$




Classification Problem



Matrix Multiplication


$$\begin{bmatrix} 1 & x_1 & x_2 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} = X^T W$$

$(1, 3)$ $(3, 1)$ $(1, 1)$





$$\begin{bmatrix} 1 & x_1^1 & x_2^1 \\ 1 & x_1^2 & x_2^2 \\ 1 & x_1^3 & x_2^3 \\ 1 & x_1^4 & x_2^4 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} = X^T W$$

$(4, 3)$ $(3, 1)$ $(4, 1)$





Multilayer Perceptron




Matrix Multiplication

$$[x_1 \quad x_2] \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} \begin{bmatrix} w_{11}^{(2)} \\ w_{21}^{(2)} \end{bmatrix} = XW^{(1)}W^{(2)}$$


$(1, 2)$ $(2, 2)$ $(2, 1)$
 Equal! Equal!



Artificial Neural Network




Artificial Neural Network (completed)



Softmax

$$\text{softmax}(x_i) = \frac{\exp(x_i)}{\sum_{j=0}^n \exp(x_j)}$$



03 How to Train A Model




Procedure of Training a DL Model


- ▲ Step 1 | Arrange data into in “tensor” format
- ▲ Step 2 | Define the network
- ▲ Step 3 | Define a loss function and an optimizer
- ▲ Step 4 | Train the model
- ▲ Step 5 | Test the model

MNIST Dataset

Handwritten digits



Define a Model in PyTorch



Error Rate (1-Accuracy)

Model 1 Error rate = 0.25

		Dog	Cat	Snake
Image 1	Cat	0.3	0.4	0.3
Image 2	Dog	0.5	0.3	0.2
Image 3	Cat	0.6	0.1	0.3
Image 4	Snake	0.3	0.2	0.5


Model 2 Error rate = 0.25

		Dog	Cat	Snake
Image 1	Cat	0.1	0.8	0.1
Image 2	Dog	0.9	0.1	0.0
Image 3	Cat	0.4	0.3	0.3
Image 4	Snake	0.2	0.0	0.8


Mean Square Error

$$MSE = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

Model 1 **MSE = 0.18**



Model 2 **MSE = 0.02**




One hot encoding


Cross-Entropy

$$CE = \sum_{c=1}^C \sum_{i=1}^n -y_{c,i} \log_2(p_{c,i})$$

Model 1 **CE = 0.22**



Model 2 **CE = 0.09**




Solving the Loss Function


Typically we estimate a parameter by minimizing the loss function, and using as the estimator the parameter which minimizes the loss.

$$\min_w \text{loss}(w)$$

Usually (but not always) the way to solve the loss function is to differentiate it and equate it to zero.



Gradient Descent



Pick a initial value w^0 (randomly)



While True:

$$w^{t+1} \leftarrow w^t - \eta \frac{\partial f}{\partial w} \Big|_{w=w^t}$$

Learning Rate
(hyperparameter)

Gradient

Optimizers






04 Convolutional Neural Networks


Challenges with ANN

- ANN loses the spatial features of an image.
- The number of trainable parameters increases drastically with an increase in the size of the image.




Convolutional Neural Networks

Convolution + Artificial Neural Network




Convolution (a review)



Laplacian Filter


0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1



CNN learns the filters by itself

Convolution Layer



Max Pooling

Max pooling is done by applying a max filter to (usually) non-overlapping subregions of the initial representation.

146	155	144	130	145	151
142	153	150	128	131	151
131	141	142	130	128	148
122	123	125	127	130	135
130	123	107	118	150	154
127	120	125	143	153	161



155	150	151
141	142	148
130	143	161


It reduces the computational cost by reducing the number of parameters to learn and provides basic translation invariance to the internal representation.

This layer requires no training!

Convolutional Layer in PyTorch


```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```



Convolutional Layer in PyTorch

```
def forward(self, x):  
    x = self.pool(F.relu(self.conv1(x)))  
    x = self.pool(F.relu(self.conv2(x)))  
    x = x.view(-1, 16 * 5 * 5)  
    x = F.relu(self.fc1(x))  
    x = F.relu(self.fc2(x))  
    x = self.fc3(x)  
    return x
```

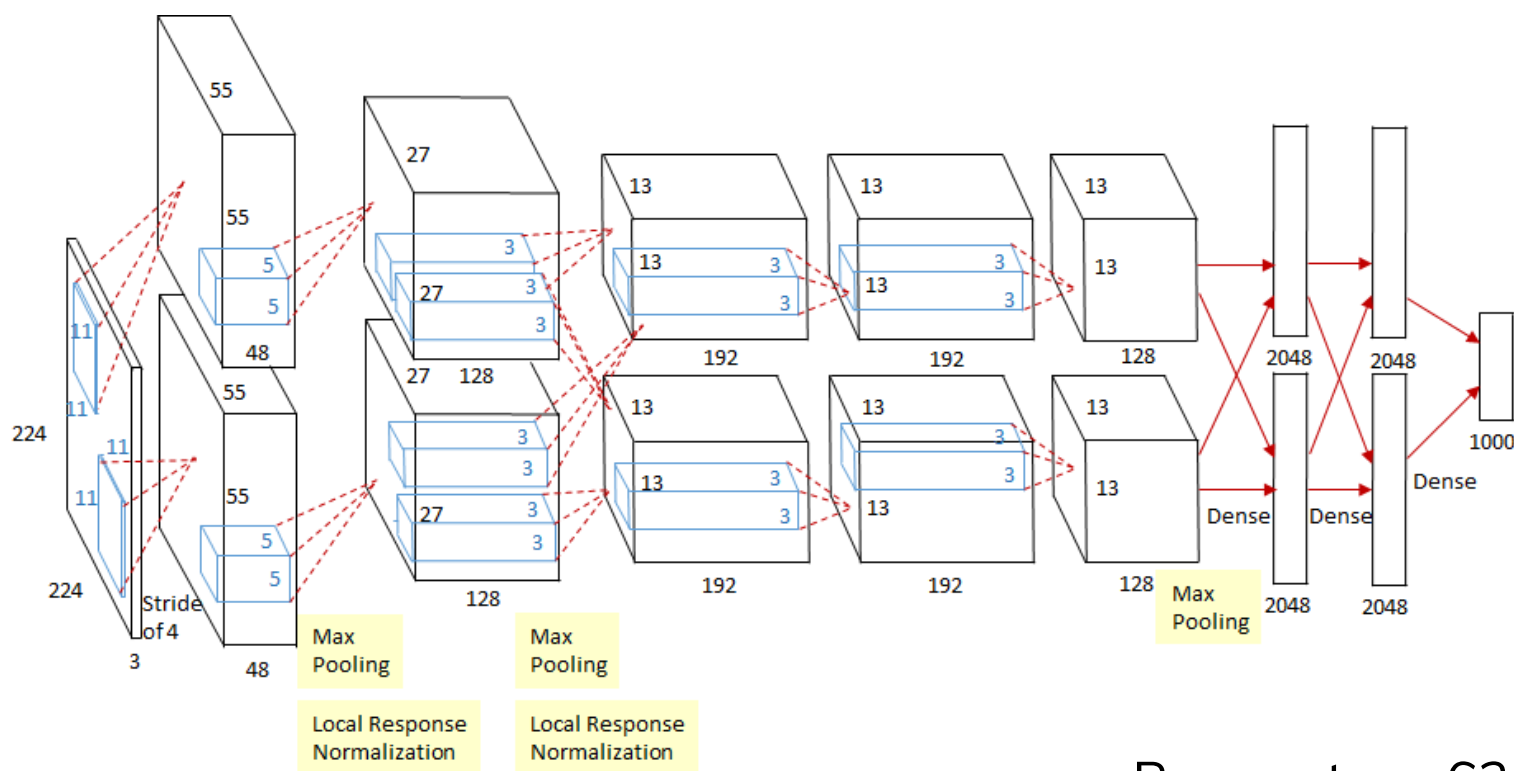


torch.nn

Deep neural network (DNN) models can be composed just like building LEGO buildings




CNN Architectures | AlexNet



Parameters: 62M



Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. In Advances in neural information processing systems (pp. 1097-1105).

CNN Architectures | VGG

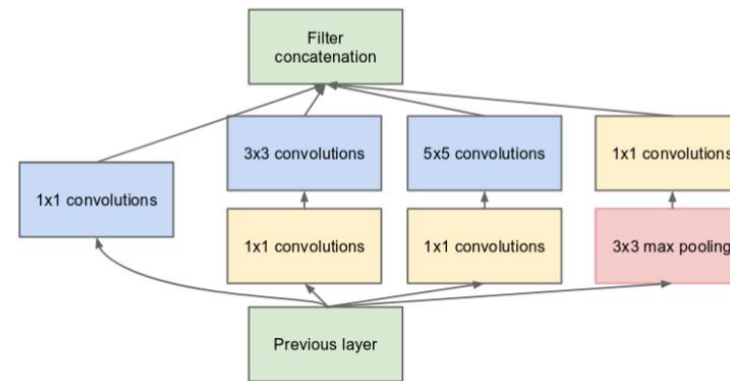


Parameters: 138M

CNN Architectures | GoogLeNet


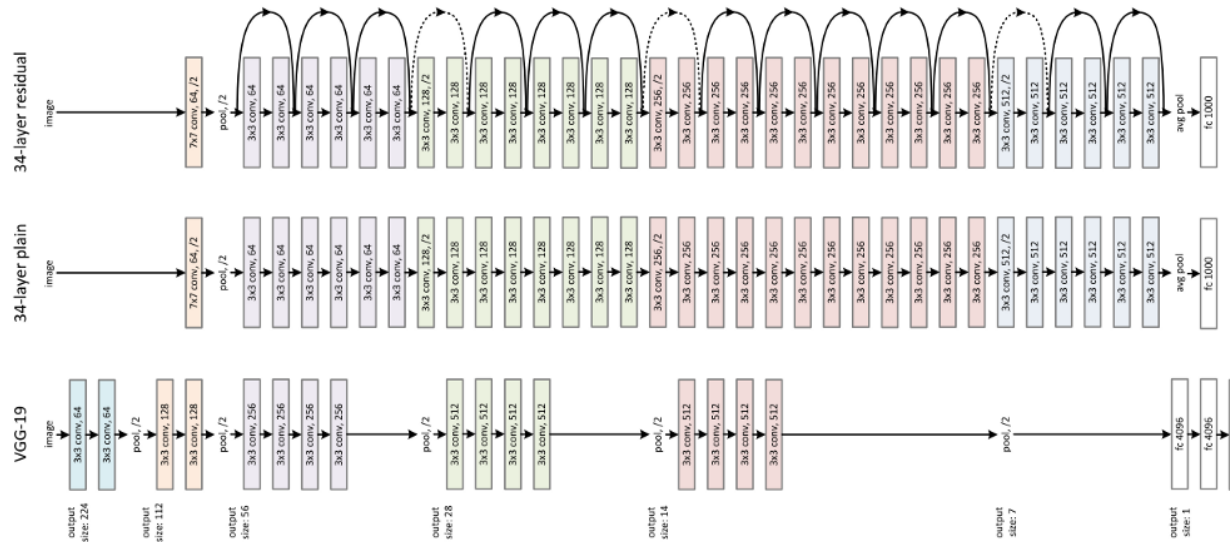


Inception module




Parameters: 4M

CNN Architectures | ResNet



Parameters: 11-58M

CNN Architectures | Summery



Computer Vision Tasks

Semantic Segmentation Classification + Localization




TRUNK, CAT, LEAF

Only pixels, No object

CAT

Single Object

Object Detection



CAT, DOG, DOG, CAT


Multiple Objects

Instance Segmentation



CAT, DOG, DOG, CAT

Deep Learning Frameworks



thank you!